

The R Cookbook

Jeffrey Wong

November 16, 2012

Contents

1	When to Use R?	2
2	The Basic Operations	2
2.1	Data Structures	2
3	Getting Data into R	5
3.1	Accessing Data	6
4	Data Munging by Example	10
4.1	The <i>apply</i> Family	10
4.2	String manipulation	12
4.3	Reshaping - Handling Panel Data	13
5	Statistics Functions	15
6	Data Visualization	16
6.1	Histograms and PDFs	17
6.2	Line Graphs, Smoothing	19
6.3	Faceting the Electric Load Data	21
7	Data Modeling	25
7.1	Regression	26
7.2	Generalized Linear Models	26
7.3	Neural Networks	26
7.4	CART: Classification and Regression Tree	26
8	Advanced Data Analysis	26
8.1	Penalized Models	26
8.2	Boosting	26
9	Open Source Development	26
9.1	devtools	26
9.2	Roxygen	26
9.3	Non-R Source Code	26

1 When to Use R?

R is the go-to language for statisticians; it's free, open source, contains a lot of statistics functions and is rapidly being developed. The R community has generated thousands of packages that serve as add-ons; the rising popularity in statistical learning has also accelerated growth in the R community. While MATLAB provides a toolbox of matrix algebra and other types of analysis derived from matrix algebra, R provides extensive services in probability, a statistician's best friend. Going beyond MATLAB, R can operate on both categorical and numerical data, making R a natural tool for social scientists while MATLAB is the dominant tool for engineers and physical scientists. Since R is easily distributable, it is popular both in the classroom and in industry.

Anything you can do in MATLAB, you can do in R. Anything you can do in scikits you can do in R. However, that does not mean R should be a dominant language. R has one of the steepest learning curves, mainly because of its poor documentation. Also, R is not a formal programming language, so it lacks the ability to operate in production environments. For example, error handling is no where near as elegant as in standard programming languages, such as Java and Python. R has developed a culture of being only used in a development environment - to experiment with. For most statistical analysis, this is enough; if you are performing one study R might very well be your best bet. If you are a data company however, you might use R to explore data but might write your own functions in a traditional production language like Java.

In summary, R has all the tools that a statistician would ever need. However, since R is not backed by any company and is only developed by a pool of volunteers, it has not established a culture of production level code. If you are just looking to gain insight into your own data, R is the best, but if you are building an analytics service to understand other people's data, you should consider another language.

My own style is to use R for data exploration. If speed is necessary, I will switch to scikits or MATLAB. For production, I either use Java libraries such as Weka or Mahout, or Python's numpy & scipy libraries.

2 The Basic Operations

2.1 Data Structures

There are three basic data structures in R that everyone needs to know.

- Vector - a 1D array. The data types in a vector must all be the same
- List - R's most generic container. List holds elements that are not necessarily the same data type
- Matrix - a data structure with standard matrix operations. Can contain numeric data only

- Data Frame - a 2D table that can contain both numeric and categorical data

A vector is created by concatenating elements together

```
> x = c(1,2,3,4,5)
> x[1]
[1] 1
> x[5]
[1] 5
> z = c("HELLO", "WORLD")
> z
[1] "HELLO" "WORLD"
```

A list can contain any amount of named, or unnamed components. The components can be any data structure, even another list. To access the components of a list, we use the `$` operator; in most programming languages, the `.` is used

```
> x = list(x = c(1,2,3,4,5),
+           y = c(6,7,8,9,10),
+           z = c("HELLO", "WORLD"),
+           list("HOW ARE", "YOU?"))
> x$x
[1] 1 2 3 4 5
> x$y
[1] 6 7 8 9 10
> x$z
[1] "HELLO" "WORLD"
> x[[1]]
[1] 1 2 3 4 5
> x[[2]]
[1] 6 7 8 9 10
> x[[3]]
[1] "HELLO" "WORLD"
```

```

> x[[4]]
[[1]]
[1] "HOW ARE"
[[2]]
[1] "YOU?"

```

Data frames are really a list of vectors. You can even access elements of a data frame the same way you access a list, through the `$` operator!

```

> x = data.frame(variable1 = c(1,2,3,4,5),
+                  variable2 = c(6,7,8,9,10))
> x

  variable1 variable2
1          1          6
2          2          7
3          3          8
4          4          9
5          5         10

> x$variable1

[1] 1 2 3 4 5

```

One useful way to generate random data for your algorithms is to use the `rnorm` function. `rnorm(n)`, without any other arguments will generate n random $N(0, 1)$ variables.

```

> randomData = rnorm(100)
> x = matrix(randomData, nrow = 10, ncol = 10)
> x

 [,1]      [,2]      [,3]      [,4]      [,5]      [,6]
[1,] -0.5666295 -0.4537955  0.9238866 -0.2282495 -0.58781614  1.3404770
[2,] -0.2607414  0.1364072 -0.7752184  0.5623043  1.21158796  1.0504152
[3,]  1.7447843 -1.5910087 -1.3105325  0.2097337 -0.03616673  0.6225778
[4,]  0.3829827  0.4790765  0.9155848 -0.8451527  0.97099090 -0.7658583
[5,] -0.4666524  0.1638046 -1.5233620  0.6377669 -1.55283407 -0.7624937
[6,]  0.6864521  0.1260894 -0.3944970  0.4320908 -0.21518549  0.5621253
[7,] -1.4610579 -0.4189762  0.6132362  0.5400851 -0.64227816 -1.2079022
[8,] -0.3810007  1.4856098  0.5613989 -0.5537582  2.50517633  0.4948815
[9,] -0.5336655  0.8617824 -0.2835345  0.5470728  1.12824087  0.3668310
[10,] 1.0083807 -0.6919587 -0.7771089 -0.7605939 -0.41303793 -1.1409753
                [,7]      [,8]      [,9]      [,10]
[1,]  0.43831345  1.09714081  0.6022230  1.96805914
[2,]  1.36257898 -1.21628591 -1.3741138  0.94328092

```

```
[3,] -1.26718549  0.93211424  0.8728925  0.28863163
[4,]  0.29765699 -0.36234966 -0.9987472 -0.43087156
[5,]  0.42017192  0.02094912  2.0062024  0.61818046
[6,] -0.12773294  0.01993031 -0.6098149  0.05286031
[7,]  2.43427989 -1.00025381 -0.6771333  0.70387698
[8,]  0.03361074 -1.48664261  0.5622671 -1.14210028
[9,] -0.65635741 -0.30073922 -2.1218081  0.18493460
[10,] 0.48107895  0.57764446  0.4909524 -1.21139020
```

The function *matrix* creates a 10 by 10 matrix and fills it with *randomData*. By default, it takes the elements from the 1D vector and fills the matrix by columns, that is the first 10 elements of *randomData* form the first column of *x*, then the next 10 elements form the second column etc. To fill it by row, we can use the operational parameter *byrow*

```
> x = matrix(randomData, nrow = 10, ncol = 10, byrow = TRUE)
> x

[,1]      [,2]      [,3]      [,4]      [,5]      [,6]
[1,] -0.5666295 -0.2607414  1.74478434  0.3829827 -0.46665238  0.68645211
[2,] -0.4537955  0.1364072 -1.59100873  0.4790765  0.16380462  0.12608939
[3,]  0.9238866 -0.7752184 -1.31053254  0.9155848 -1.52336204 -0.39449703
[4,] -0.2282495  0.5623043  0.20973372 -0.8451527  0.63776694  0.43209077
[5,] -0.5878161  1.2115880 -0.03616673  0.9709909 -1.55283407 -0.21518549
[6,]  1.3404770  1.0504152  0.62257784 -0.7658583 -0.76249365  0.56212526
[7,]  0.4383134  1.3625790 -1.26718549  0.2976570  0.42017192 -0.12773294
[8,]  1.0971408 -1.2162859  0.93211424 -0.3623497  0.02094912  0.01993031
[9,]  0.6022230 -1.3741138  0.87289249 -0.9987472  2.00620238 -0.60981493
[10,] 1.9680591  0.9432809  0.28863163 -0.4308716  0.61818046  0.05286031

[,7]      [,8]      [,9]      [,10]
[1,] -1.4610579 -0.38100072 -0.5336655  1.0083807
[2,] -0.4189762  1.48560983  0.8617824 -0.6919587
[3,]  0.6132362  0.56139888 -0.2835345 -0.7771089
[4,]  0.5400851 -0.55375821  0.5470728 -0.7605939
[5,] -0.6422782  2.50517633  1.1282409 -0.4130379
[6,] -1.2079022  0.49488150  0.3668310 -1.1409753
[7,]  2.4342799  0.03361074 -0.6563574  0.4810790
[8,] -1.0002538 -1.48664261 -0.3007392  0.5776445
[9,] -0.6771333  0.56226711 -2.1218081  0.4909524
[10,] 0.7038770 -1.14210028  0.1849346 -1.2113902
```

3 Getting Data into R

R reads data that is in a table format. This includes any delimited file type like csv, and tsv. R also handles spreadsheet formats from excel, as well as SAS

data files. The most common function you will use is `read.csv`, which is a bit misleading because it can handle any delimited file type, not just csv.

In this cookbook, I have included a dataset from Kaggle on electric loads. In this dataset, there are 20 geographic zones, with measurements of electric loads taken every hour. If your current working directory is inside this project folder, then you can issue this command to read the data.

```
> data = read.csv("data/Load_history.csv", header=T)
```

You can even open data files that are hosted online

```
> #data = read.csv("https://raw.github.com/", header=T)
```

3.1 Accessing Data

Most basic statistics operations can be performed with just one command in R.

Now that we have data, we can take a peek at what it looks like. Use `head` and `tail` to peek at a few rows.

```
> head(data)
```

zone_id	year	month	day	h1	h2	h3	h4	h5	h6	h7				
1	1	2004	1	16,853	16,450	16,517	16,873	17,064	17,727	18,574				
2	1	2004	1	2	14,155	14,038	14,019	14,489	14,920	16,072	17,800			
3	1	2004	1	3	14,439	14,272	14,109	14,081	14,775	15,491	16,536			
4	1	2004	1	4	11,273	10,415	9,943	9,859	9,881	10,248	11,016			
5	1	2004	1	5	10,750	10,321	10,107	10,065	10,419	12,101	14,847			
6	1	2004	1	6	15,742	15,682	16,132	16,761	17,909	20,234	23,948			
				h8	h9	h10	h11	h12	h13	h14	h15	h16	h17	h18
1	19,355	19,534	18,611	17,666	16,374	15,106	14,455	13,518	13,138	14,130	16,809			
2	19,089	19,577	20,047	19,770	18,564	18,137	17,046	16,127	15,448	15,839	17,727			
3	18,197	19,109	18,012	17,200	15,950	14,978	14,162	13,507	13,414	13,826	15,825			
4	12,780	15,108	15,680	15,280	14,605	14,689	14,642	14,207	13,614	14,162	16,237			
5	15,259	14,045	14,009	14,332	13,908	13,981	13,865	13,845	14,350	15,501	17,307			
6	24,789	23,024	20,613	19,070	17,447	17,556	18,506	18,762	19,162	21,509	25,314			
				h19	h20	h21	h22	h23	h24					
1	18,150	18,235	17,925	16,904	16,162	14,750								
2	18,895	18,650	18,443	17,580	16,467	15,258								
3	16,996	16,394	15,406	14,278	13,315	12,424								
4	17,430	17,218	16,633	15,238	13,580	11,727								
5	18,786	19,089	19,192	18,416	17,006	16,018								
6	28,060	28,768	28,919	28,653	27,406	26,507								

```
> tail(data)
```

zone_id year month day h1 h2 h3 h4 h5 h6 h7 h8 h9 h10 h11 h12 h13 h14 h15
32995 20 2008 7 2

```

32996      20 2008      7   3
32997      20 2008      7   4
32998      20 2008      7   5
32999      20 2008      7   6
33000      20 2008      7   7
               h16 h17 h18 h19 h20 h21 h22 h23 h24
32995
32996
32997
32998
32999
33000

```

Here you can see what data types you have, in this case all numerics. Also, you can see that there is a lot of missing data at the bottom of the file.

To subset data by taking the first 10 rows

```
> data[1:10,]
```

	zone_id	year	month	day	h1	h2	h3	h4	h5	h6	h7				
1	1	2004	1	1	16,853	16,450	16,517	16,873	17,064	17,727	18,574				
2	1	2004	1	2	14,155	14,038	14,019	14,489	14,920	16,072	17,800				
3	1	2004	1	3	14,439	14,272	14,109	14,081	14,775	15,491	16,536				
4	1	2004	1	4	11,273	10,415	9,943	9,859	9,881	10,248	11,016				
5	1	2004	1	5	10,750	10,321	10,107	10,065	10,419	12,101	14,847				
6	1	2004	1	6	15,742	15,682	16,132	16,761	17,909	20,234	23,948				
7	1	2004	1	7	26,014	26,447	27,286	27,923	29,130	31,503	34,900				
8	1	2004	1	8	25,104	25,122	25,464	25,715	26,219	28,552	31,815				
9	1	2004	1	9	21,175	21,056	21,241	22,062	23,026	25,610	27,220				
10	1	2004	1	10	23,405	23,507	24,067	24,786	25,418	26,631	28,560				
					h8	h9	h10	h11	h12	h13	h14	h15	h16	h17	h18
1	19,355	19,534	18,611	17,666	16,374	15,106	14,455	13,518	13,138	14,130	16,809				
2	19,089	19,577	20,047	19,770	18,564	18,137	17,046	16,127	15,448	15,839	17,727				
3	18,197	19,109	18,012	17,200	15,950	14,978	14,162	13,507	13,414	13,826	15,825				
4	12,780	15,108	15,680	15,280	14,605	14,689	14,642	14,207	13,614	14,162	16,237				
5	15,259	14,045	14,009	14,332	13,908	13,981	13,865	13,845	14,350	15,501	17,307				
6	24,789	23,024	20,613	19,070	17,447	17,556	18,506	18,762	19,162	21,509	25,314				
7	35,201	32,405	29,694	27,298	25,243	23,481	22,095	20,617	21,013	23,676	27,329				
8	32,289	28,968	25,705	23,297	23,090	23,244	23,081	22,421	22,883	24,436	26,555				
9	27,570	27,422	26,881	26,574	25,648	24,980	24,682	24,955	24,932	25,497	27,668				
10	30,242	32,387	31,926	29,308	27,017	24,909	23,193	22,257	22,457	23,909	27,515				
					h19	h20	h21	h22	h23	h24					
1	18,150	18,235	17,925	16,904	16,162	14,750									
2	18,895	18,650	18,443	17,580	16,467	15,258									
3	16,996	16,394	15,406	14,278	13,315	12,424									
4	17,430	17,218	16,633	15,238	13,580	11,727									
5	18,786	19,089	19,192	18,416	17,006	16,018									

```

6 28,060 28,768 28,919 28,653 27,406 26,507
7 29,685 29,838 29,806 28,704 27,069 25,708
8 27,394 27,486 26,890 25,529 23,869 22,278
9 28,784 28,113 27,311 26,327 24,967 23,824
10 29,526 30,073 30,858 30,698 30,208 30,056

```

and the first 10 columns

```

> first10 = data[,1:10]
> head(first10)

```

	zone_id	year	month	day	h1	h2	h3	h4	h5	h6
1	1	2004	1	1	16,853	16,450	16,517	16,873	17,064	17,727
2	1	2004	1	2	14,155	14,038	14,019	14,489	14,920	16,072
3	1	2004	1	3	14,439	14,272	14,109	14,081	14,775	15,491
4	1	2004	1	4	11,273	10,415	9,943	9,859	9,881	10,248
5	1	2004	1	5	10,750	10,321	10,107	10,065	10,419	12,101
6	1	2004	1	6	15,742	15,682	16,132	16,761	17,909	20,234

To see the names of the columns

```
> names(data)
```

```

[1] "zone_id"   "year"      "month"     "day"       "h1"        "h2"        "h3"
[8] "h4"         "h5"        "h6"        "h7"        "h8"        "h9"        "h10"
[15] "h11"        "h12"        "h13"        "h14"        "h15"        "h16"        "h17"
[22] "h18"        "h19"        "h20"        "h21"        "h22"        "h23"        "h24"

```

To select the year column by its name

```
> head(data$year)
```

```
[1] 2004 2004 2004 2004 2004 2004
```

To find something in a vector, we can use the *which* command. The output of which tells us the location of an object within the vector. In this example, we will look at the month column of the data frame, which on its own is a vector.

```
> head(which(data$month == 1))
```

```
[1] 1 2 3 4 5 6
```

To select the rows of the data frame that belong to January, we can apply *which* to the subsetting of rows

```

> january = data[which(data$month == 1),]
> head(january)

```

	zone_id	year	month	day	h1	h2	h3	h4	h5	h6	h7			
1	1	2004	1	1	16,853	16,450	16,517	16,873	17,064	17,727	18,574			
2	1	2004	1	2	14,155	14,038	14,019	14,489	14,920	16,072	17,800			
3	1	2004	1	3	14,439	14,272	14,109	14,081	14,775	15,491	16,536			
4	1	2004	1	4	11,273	10,415	9,943	9,859	9,881	10,248	11,016			
5	1	2004	1	5	10,750	10,321	10,107	10,065	10,419	12,101	14,847			
6	1	2004	1	6	15,742	15,682	16,132	16,761	17,909	20,234	23,948			
				h8	h9	h10	h11	h12	h13	h14	h15	h16	h17	h18
1	19,355	19,534	18,611	17,666	16,374	15,106	14,455	13,518	13,138	14,130	16,809			
2	19,089	19,577	20,047	19,770	18,564	18,137	17,046	16,127	15,448	15,839	17,727			
3	18,197	19,109	18,012	17,200	15,950	14,978	14,162	13,507	13,414	13,826	15,825			
4	12,780	15,108	15,680	15,280	14,605	14,689	14,642	14,207	13,614	14,162	16,237			
5	15,259	14,045	14,009	14,332	13,908	13,981	13,865	13,845	14,350	15,501	17,307			
6	24,789	23,024	20,613	19,070	17,447	17,556	18,506	18,762	19,162	21,509	25,314			
				h19	h20	h21	h22	h23	h24					
1	18,150	18,235	17,925	16,904	16,162	14,750								
2	18,895	18,650	18,443	17,580	16,467	15,258								
3	16,996	16,394	15,406	14,278	13,315	12,424								
4	17,430	17,218	16,633	15,238	13,580	11,727								
5	18,786	19,089	19,192	18,416	17,006	16,018								
6	28,060	28,768	28,919	28,653	27,406	26,507								

To see the data structures behind the object *data*.

```
> str(data)

'data.frame':      33000 obs. of  28 variables:
 $ zone_id: int  1 1 1 1 1 1 1 1 1 ...
 $ year   : int  2004 2004 2004 2004 2004 2004 2004 2004 2004 ...
 $ month  : int  1 1 1 1 1 1 1 1 1 ...
 $ day    : int  1 2 3 4 5 6 7 8 9 10 ...
 $ h1     : Factor w/ 25095 levels "", "100,009", "100,016", ...: 8072 4877 5240 999 517 6843 137 ...
 $ h2     : Factor w/ 24957 levels "", "10,001", "100,017", ...: 8099 5331 5621 401 292 7293 137 ...
 $ h3     : Factor w/ 24946 levels "", "1", "10,000", ...: 8273 5551 5663 24902 110 7904 13894 137 ...
 $ h4     : Factor w/ 25037 levels "", "10,000", "100,028", ...: 8522 6203 5678 24911 59 8435 14 ...
 $ h5     : Factor w/ 25067 levels "", "10,000", "10,001", ...: 8606 6580 6434 24957 418 9211 14 ...
 $ h6     : Factor w/ 25317 levels "", "100,001", "100,007", ...: 8745 7288 6621 230 2046 10400 ...
 $ h7     : Factor w/ 25683 levels "", "10,000", "100,002", ...: 8979 8201 6713 1034 4495 12560 ...
 $ h8     : Factor w/ 25713 levels "", "100,028", "10,003", ...: 9298 9010 8038 2533 4552 12910 ...
 $ h9     : Factor w/ 25654 levels "", "100,009", "100,010", ...: 9163 9206 8708 3965 3041 11799 ...
 $ h10    : Factor w/ 25629 levels "", "10,000", "100,011", ...: 7700 9266 6935 4160 2838 9836 1 ...
 $ h11    : Factor w/ 25603 levels "", "100,003", "100,011", ...: 6304 8764 5696 3680 2982 8004 ...
 $ h12    : Factor w/ 25548 levels "", "100,003", "100,020", ...: 4793 7561 4373 3232 2708 6141 ...
 $ h13    : Factor w/ 25591 levels "", "100,010", "100,019", ...: 3786 7202 3672 3429 2825 6458 ...
 $ h14    : Factor w/ 25626 levels "", "10,001", "100,022", ...: 3304 6001 3046 3473 2778 7794 1 ...
 $ h15    : Factor w/ 25558 levels "", "0", "100,000", ...: 2546 4983 2542 3138 2822 8155 9852 1 ...
 $ h16    : Factor w/ 25631 levels "", "0", "10,000", ...: 2249 4235 2475 2638 3268 8517 10116 1 ...
```

```

$ h17 : Factor w/ 25732 levels "", "10,000", "100,012", ... : 3057 4588 2818 3090 4261 10475
$ h18 : Factor w/ 25947 levels "", "1", "100,005", ... : 5473 6522 4491 4886 6031 13137 14252
$ h19 : Factor w/ 26046 levels "", "0", "100,019", ... : 6703 7513 5481 5949 7400 14597 15242
$ h20 : Factor w/ 25936 levels "", "100,005", "100,016", ... : 6606 7086 4815 5563 7572 14907
$ h21 : Factor w/ 25821 levels "", "10,000", "100,009", ... : 6180 6819 3925 4850 7729 14957
$ h22 : Factor w/ 25697 levels "", "100,004", "10,001", ... : 5301 6146 3186 3779 7200 14705
$ h23 : Factor w/ 25450 levels "", "10,000", "100,001", ... : 5083 5501 2343 2503 6263 13953
$ h24 : Factor w/ 25200 levels "", "100,010", "100,016", ... : 4587 5250 1571 1089 6183 13571

```

That's interesting, what are all these factors? Factors are a data type that R uses to represent categorical variables. Here, *h1* is a factor with 25,095 levels, meaning *h1* is a category and it can fall under one of 25,095 different values.

4 Data Munging by Example

Keywords: data manipulation, cleaning

We know that energy usage should not be a categorical variable. Why did R detect them as such? In this particular file, it is because of the "," in the numerics. Spreadsheet programs might display ","'s in their UI, but typically there is no literal "," in the number. When the text file was created, the "," was outputted and is now confusing R.

Let's use some regexes to strip the comma out. First, I'll introduce a set of functions that are handy to have when manipulating data structures. Then, we will move into manipulating the text of this data file.

4.1 The *apply* Family

Keyword: map, apply

In functional programming, there is always a function called *map*, which applies a function to each element of an array. Say we had the array $x = [1,2,3,4,5]$ and defined a function $f(x) = x^2$. Then $x.map(f)$ would output $[1, 4, 9, 16, 25]$. In R, there is a family of such functions called the *apply* family. The R function *sapply* operates on R vectors and does exactly what *map* does.

```

> x = c(1,2,3,4,5)
> sapply(x, function(i) { return (i^2) } )
[1]  1  4  9 16 25

```

Here, the parameter *i* that gets passed in to the function that we define is an element from *x*.

The function *apply* goes across the rows (or columns) of either a data frame or matrix and applies a function. Say we wanted to double every row of a matrix *x*

```

> x = matrix(rnorm(100), 10, 10)
> apply(x, 1, function(i) { return (2*i) } )

```

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]
[1,]	-0.9214741	-0.4080685	0.6178581	0.2526408	0.2930994	0.1995567
[2,]	0.1582797	-0.3386092	-1.8394219	0.1374866	0.9578987	-4.1670646
[3,]	-1.0988274	0.1874387	0.5990334	1.7749538	-0.3818964	1.8480539
[4,]	3.0684138	-2.5156762	2.1459080	3.1692040	-2.9252367	0.3070902
[5,]	3.0430428	-2.4078486	0.9365027	2.0644205	3.2691936	-1.9842970
[6,]	-3.3392650	2.8545964	1.0777143	-1.8194098	2.1310400	1.9301158
[7,]	1.5651559	0.1401914	-3.3306147	-0.2818149	-1.4541628	-0.6087199
[8,]	-0.8364001	-0.6231714	-2.0163279	1.4030175	-1.7367937	-0.7989320
[9,]	-4.9457231	-1.6224378	-2.5488993	-1.7109441	1.4727127	-0.4898371
[10,]	-1.9455697	-0.5896207	2.0737743	-0.7803228	2.0163871	-0.6444963
	[,7]	[,8]	[,9]	[,10]		
[1,]	0.08461708	1.4628028	-1.0517544	-3.6204459		
[2,]	2.04048989	-2.3913849	-1.9423340	2.6614654		
[3,]	-0.51883520	1.6546440	1.2014253	-0.3905317		
[4,]	0.86305411	0.2649672	-2.0139563	2.8121050		
[5,]	0.11897523	-0.1404777	-3.7866794	0.5859579		
[6,]	6.21078115	0.9627342	-0.3303900	0.4833403		
[7,]	1.54605919	0.2576065	0.1139882	-0.3150234		
[8,]	1.09560821	-0.2852498	-1.0997969	-1.1833216		
[9,]	1.54793054	0.4152166	-1.7724261	0.5280253		
[10,]	-1.81942744	-0.3043483	2.2575883	3.2396727		

If we want to double the columns

```
> x = matrix(rnorm(100), 10, 10)
> apply(x, 2, function(j) { return (2*j) } )



|       | [,1]        | [,2]        | [,3]        | [,4]       | [,5]       | [,6]       |
|-------|-------------|-------------|-------------|------------|------------|------------|
| [1,]  | 1.32676833  | 1.56259647  | 4.0289326   | 1.1892541  | -0.1955158 | 1.5907255  |
| [2,]  | -1.53225344 | 3.21129913  | 1.0993994   | 2.1452932  | -0.3872460 | 0.5366756  |
| [3,]  | 0.92788352  | 3.36677617  | 3.1427030   | 0.9207192  | -0.1436585 | -0.7583696 |
| [4,]  | 0.53218777  | 0.07661448  | 1.6543962   | -1.0547685 | -2.2549091 | 1.5301064  |
| [5,]  | 0.30484101  | 1.11083747  | 1.4357998   | -1.4887626 | -2.8318896 | -2.1727493 |
| [6,]  | 0.04519007  | -3.28781963 | -1.4144825  | -2.3803700 | -1.4612739 | 0.6269776  |
| [7,]  | 1.59585503  | 3.38780458  | 2.0168684   | 0.9403496  | -3.4821498 | -3.9100354 |
| [8,]  | 0.73829648  | 1.73990504  | 0.3099872   | -3.3903698 | 0.3623023  | -1.7288807 |
| [9,]  | -1.40070554 | 3.71508668  | -1.9184296  | -2.9428822 | -1.8188117 | 1.8945809  |
| [10,] | 0.51898448  | 4.83692321  | 1.1912197   | -0.6928451 | -3.6855648 | -3.2822222 |
|       | [,7]        | [,8]        | [,9]        | [,10]      |            |            |
| [1,]  | -0.61947572 | -1.3605925  | -0.62045315 | -0.1725707 |            |            |
| [2,]  | -1.13956171 | 0.3105300   | -1.99974412 | 2.7628658  |            |            |
| [3,]  | -2.10172551 | -0.1012705  | 2.76732054  | 0.4758645  |            |            |
| [4,]  | 1.11836885  | 1.2532202   | 0.04572582  | -3.7348241 |            |            |
| [5,]  | 0.03773514  | -4.6293148  | 1.05585305  | 3.2945779  |            |            |
| [6,]  | 0.31285424  | -3.7803802  | 1.25776586  | -0.2978130 |            |            |
| [7,]  | 2.26260056  | 0.3784190   | 0.34382279  | -3.1547990 |            |            |


```

```
[8,]  1.35597388 -0.9450148 -0.53882869  0.1426317
[9,] -0.81009810  0.2462061  0.12198392  2.8445674
[10,] -0.19466032  1.8683834  0.17579215  1.4009906
```

4.2 String manipulation

Keywords: regex

String manipulation is best done using Hadley Wickham's package *stringr*. You can install and load additional R packages by

```
> install.packages('stringr', repos = "http://cran.cnr.Berkeley.edu")
The downloaded source packages are in
  /tmp/RtmpPPwTr/downloaded_packages
```

```
> require(stringr)
```

The function we would like to use is **str_replace**. In this use case, we want to replace all commas with empty strings. Given a string with commas x, we want to run the function

```
str_replace(x, ",", "")
```

Since columns 5 through 28 all have the bad comma, we can call apply with the str_replace function.

```
> nocommas = apply(data[,5:28], 2, function(x) { str_replace_all(x, ",", "") })
> str(nocommas)

chr [1:33000, 1:24] "16853" "14155" "14439" "11273" "10750" ...
- attr(*, "dimnames")=List of 2
..$ : NULL
..$ : chr [1:24] "h1" "h2" "h3" "h4" ...
```

Note that the output of that apply gave us a data frame of all strings. To convert a string to a number, we can use the **as.numeric** function. Again, we can apply over all columns of *nocommas*

```
> loads = apply(nocommas, 2, as.numeric)
> str(loads)

num [1:33000, 1:24] 16853 14155 14439 11273 10750 ...
- attr(*, "dimnames")=List of 2
..$ : NULL
..$ : chr [1:24] "h1" "h2" "h3" "h4" ...

> head(loads)
```

```

      h1     h2     h3     h4     h5     h6     h7     h8     h9     h10    h11    h12
[1,] 16853 16450 16517 16873 17064 17727 18574 19355 19534 18611 17666 16374
[2,] 14155 14038 14019 14489 14920 16072 17800 19089 19577 20047 19770 18564
[3,] 14439 14272 14109 14081 14775 15491 16536 18197 19109 18012 17200 15950
[4,] 11273 10415 9943 9859 9881 10248 11016 12780 15108 15680 15280 14605
[5,] 10750 10321 10107 10065 10419 12101 14847 15259 14045 14009 14332 13908
[6,] 15742 15682 16132 16761 17909 20234 23948 24789 23024 20613 19070 17447
      h13    h14    h15    h16    h17    h18    h19    h20    h21    h22    h23    h24
[1,] 15106 14455 13518 13138 14130 16809 18150 18235 17925 16904 16162 14750
[2,] 18137 17046 16127 15448 15839 17727 18895 18650 18443 17580 16467 15258
[3,] 14978 14162 13507 13414 13826 15825 16996 16394 15406 14278 13315 12424
[4,] 14689 14642 14207 13614 14162 16237 17430 17218 16633 15238 13580 11727
[5,] 13981 13865 13845 14350 15501 17307 18786 19089 19192 18416 17006 16018
[6,] 17556 18506 18762 19162 21509 25314 28060 28768 28919 28653 27406 26507

```

```
> data[,5:28] = loads
```

4.3 Reshaping - Handling Panel Data

Keywords: dates, time series, panel data, reshape, melt

This dataset has components of time - year, month, and day, but they are all separated as 3 different columns. There is no actual date object here. We can create one by looping over the rows and concatenating year, month, and day into one date string; then, we will parse that string into a date object. In R, concatenation is done using the **paste** function. **paste** takes any amount of arguments, and concatenates them together with some user provided separator, i.e.

```

> paste("HELLO", "WORLD", sep=" ")
[1] "HELLO WORLD"

```

If the strings to be concatenated are stored inside a vector, then we pass in the parameter *collapse* to collapse the container into one single string

```

> words = c("HELLO", "WORLD")
> paste(words, collapse = " ")
[1] "HELLO WORLD"

```

Looping over rows lends itself to the apply function again. Each row will be considered a vector, and to construct a date string we will want to extract elements 2, 3, and 4 from the row

```

> dates = apply(data, 1, function(i) {
+   paste(i[2:4], collapse='-')
+ })
> head(dates)

```

```
[1] "2004-1-1" "2004-1-2" "2004-1-3" "2004-1-4" "2004-1-5" "2004-1-6"
```

Alternatively, paste is already vectorized, and we can call the more succinct function

```
> dates = paste(data$year, data$month, data$day, sep = "-")  
> head(dates)
```

```
[1] "2004-1-1" "2004-1-2" "2004-1-3" "2004-1-4" "2004-1-5" "2004-1-6"
```

To parse into date objects, we can use another package by Mr. Wickham, **lubridate**.

```
> install.packages('lubridate', repos="http://cran.cnr.Berkeley.edu")
```

The downloaded source packages are in
 âĂŶ/tmp/RtmpPwTr/downloaded_packagesâĂŹ

```
> require('lubridate')
```

We can parse dates in the year-month-day format using *ymd*. After parsing these date strings, we will insert the date objects into the data frame

```
> dates = ymd(dates)  
> data$date = dates  
> head(data)
```

	zone_id	year	month	day	h1	h2	h3	h4	h5	h6	h7	h8	h9	h10	h11	h12	h13	h14	h15	h16	h17	h18	h19	h20	h21	h22
1	1	2004	1	1	16853	16450	16517	16873	17064	17727	18574	19355	19534	18611	17666	16374	15106	14455	13518	13138	14130	16809	18150	18235	17925	16904
2	1	2004	1	2	14155	14038	14019	14489	14920	16072	17800	19089	19577	20047	19770	18564	18137	17046	16127	15448	15839	17727	18895	18650	18443	17580
3	1	2004	1	3	14439	14272	14109	14081	14775	15491	16536	18197	19109	18012	17200	15950	14978	14162	13507	13414	13826	15825	16996	16394	15406	14278
4	1	2004	1	4	11273	10415	9943	9859	9881	10248	11016	12780	15108	15680	15280	14605	14689	14642	14207	13614	14162	16237	17430	17218	16633	15238
5	1	2004	1	5	10750	10321	10107	10065	10419	12101	14847	15259	14045	14009	14332	13908	13981	13865	13845	14350	15501	17307	18786	19089	19192	18416
6	1	2004	1	6	15742	15682	16132	16761	17909	20234	23948	24789	23024	20613	19070	17447	17556	18506	18762	19162	21509	25314	28060	28768	28919	28653
					h23	h24																				
1	16162	14750	2004-01-01																							
2	16467	15258	2004-01-02																							
3	13315	12424	2004-01-03																							
4	13580	11727	2004-01-04																							
5	17006	16018	2004-01-05																							
6	27406	26507	2004-01-06																							

Note that we used the `$` operator to assign a non existing column, date, to the recently created `dates` vector.

This dataset is a great example of panel data. We observe 20 geographic zones over a period of 4 years. The format of this dataset has identifying variables `zone_id`, and `date`. The variable that we are observing is electric load. Each row has 24 measurements of this variable. This format, where a row stores multiple measurements on one variable is called the wide format. Sometimes, it is more helpful for a row to only contain one measurement per variable observed, called the long format. For example, how would we plot the data in the wide format? How would a plotting function know that the data moves from left to right for hours, and then from top to bottom for different dates? It would be nicer if the entire time series was captured, in order, in one column. The process of converting from wide format to long format is called **reshaping**. The package `reshape2` makes reshaping very easy, it just asks us to specify which variables are measurements.

```
> install.packages('reshape2', repos="http://cran.cnr.Berkeley.edu")

The downloaded source packages are in
  /tmp/RtmpPPwTr/downloaded_packages

> require('reshape2')
> data.long = melt(data, measure.vars = 5:28)
> head(data.long)

  zone_id year month day      date variable value
1       1 2004     1   1 2004-01-01      h1 16853
2       1 2004     1   2 2004-01-02      h1 14155
3       1 2004     1   3 2004-01-03      h1 14439
4       1 2004     1   4 2004-01-04      h1 11273
5       1 2004     1   5 2004-01-05      h1 10750
6       1 2004     1   6 2004-01-06      h1 15742
```

The long format will be very helpful in data visualization later!

5 Statistics Functions

Let's normalize the numeric parts of this data frame, contained in columns 5 through 28. We will do so by using the command `scale`, and redefining columns 5 through 28 using the output. Scale will remove the means of the columns and divide by their standard deviations

```
> data[,5:28] = scale(data[,5:28])
```

How does `scale` work behind the scenes? Note that the `scale` function can be replicated using the `apply` function. Since `scale` goes across columns, removes the mean and divides by the standard deviation, it can be written like this

```

> apply(x, 2, function(j) {
+   demean = j - mean(j)
+   return (demean / sd(j))
+ })

[,1]      [,2]      [,3]      [,4]      [,5]
[1,] 0.9817173367 -0.17554357 1.57305277 0.976519176 0.95169234
[2,] -1.7671333202 0.53138117 -0.03023195 1.477186500 0.82083039
[3,] 0.5982033713 0.59804605 1.08803405 0.835890361 0.98708663
[4,] 0.2177555803 -0.81269753 0.27350855 -0.198651179 -0.45390921
[5,] -0.0008304784 -0.36924715 0.15387441 -0.425929192 -0.84771683
[6,] -0.2504758883 -2.25528737 -1.40603763 -0.892854313 0.08777209
[7,] 1.2404349005 0.60706254 0.47188354 0.846170588 -1.29154016
[8,] 0.4159219831 -0.09951785 -0.46226406 -1.421780278 1.33242096
[9,] -1.6406545613 0.74739337 -1.68183944 -1.187435865 -0.15625888
[10,] 0.2050610761 1.22841036 0.02001977 -0.009115797 -1.43037735
          [,6]      [,7]      [,8]      [,9]      [,10]
[1,] 1.02021261 -0.49544303 -0.3277502 -0.69403791 -0.22093902
[2,] 0.52191195 -0.89700494 0.4722817 -1.78015687 1.00570781
[3,] -0.09031885 -1.63989811 0.2751367 1.97365534 0.05002612
[4,] 0.99155499 0.84635850 0.9235846 -0.16945708 -1.70951728
[5,] -0.75896490 0.01199391 -1.8926160 0.62596488 1.22789725
[6,] 0.56460210 0.22441523 -1.4861976 0.78496058 -0.27327470
[7,] -1.58026445 1.72982753 0.5047829 0.06527857 -1.46713905
[8,] -0.54912663 1.02981495 -0.1287969 -0.62976295 -0.08922369
[9,] 1.16385982 -0.64262386 0.4414874 -0.10940785 1.03984893
[10,] -1.28346664 -0.16744019 1.2180874 -0.06703670 0.43661360

```

6 Data Visualization

ggplot is the one and only way to go for data visualization in R. This visualization package is envied by many other data analysis software and by itself draws a lot of users to R. The graphs here are beautiful, easy to construct and very flexible.

All **ggplot2** commands operate on data frames, so unlike `plot(...)` you cannot pass in a vector. Fortunately, there are functions like `reshape` and `melt` that can manipulate data frames into the proper **ggplot** input. **ggplot** commands take the form `ggplot(data frame, aes(...)) + geom_something`. `aes(...)` is the aesthetics of the graph, which unintuitively means how the data points are defined: what variables make the x and y coordinates, how is the data grouped, how should the points be colored/filled/shaded. The second component, `geom_something` defines how the data is charted, whether through a line graph, a bar graph, etc. For beginners who do not know the different charting forms, there's always `qplot` which tries to infer the appropriate charting type based on the class of data in the data frame. Below are some recipes for common graphs.

6.1 Histograms and PDFs

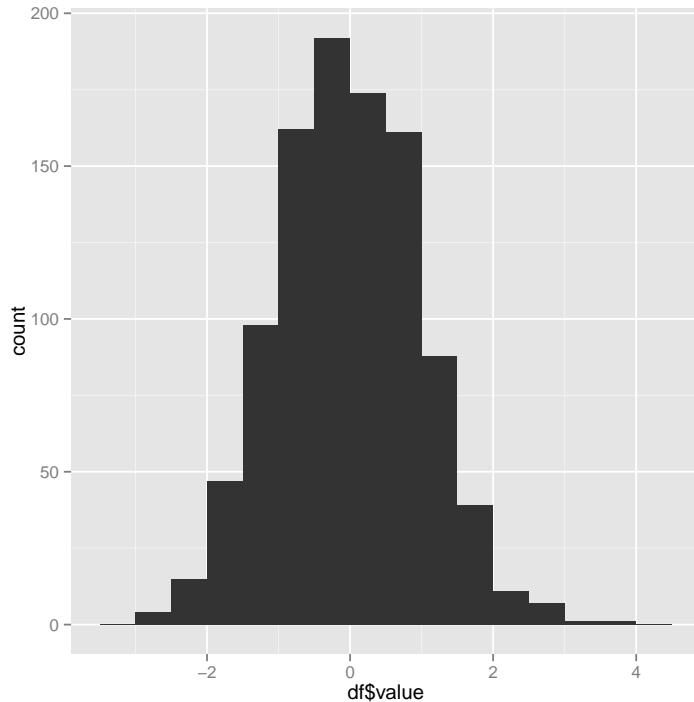
Let's create a data frame where one column is numeric, and the second column represents a group. Take a look at how qplot handles the data

```
> install.packages('ggplot2', repos="http://cran.cnr.Berkeley.edu")
```

The downloaded source packages are in
 /tmp/RtmpPwTr/downloaded_packages

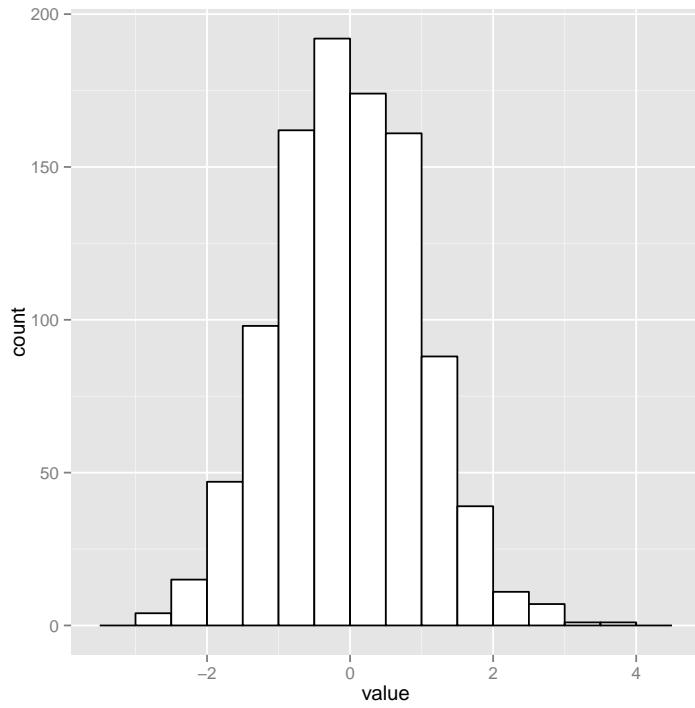
```
> require(ggplot2)
```

```
> df = data.frame(value = rnorm(1000), category = rep(factor(c("A", "B")), 500))
> qplot(df$value, binwidth=.5)
```



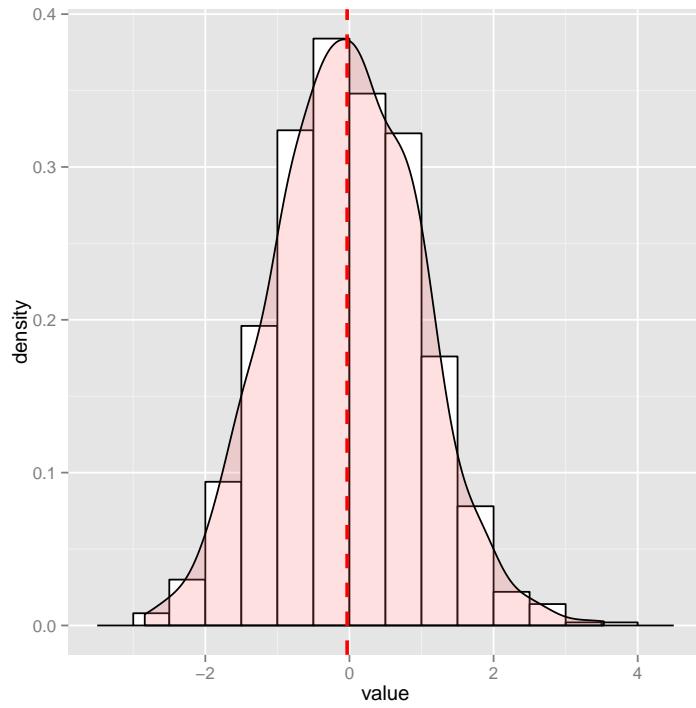
More specifically, we can control the graphing using **geom_histogram**

```
> ggplot(df, aes(x=value)) + geom_histogram(binwidth=.5, colour="black", fill="white")
```



To draw a probability density function (pdf) over the histogram

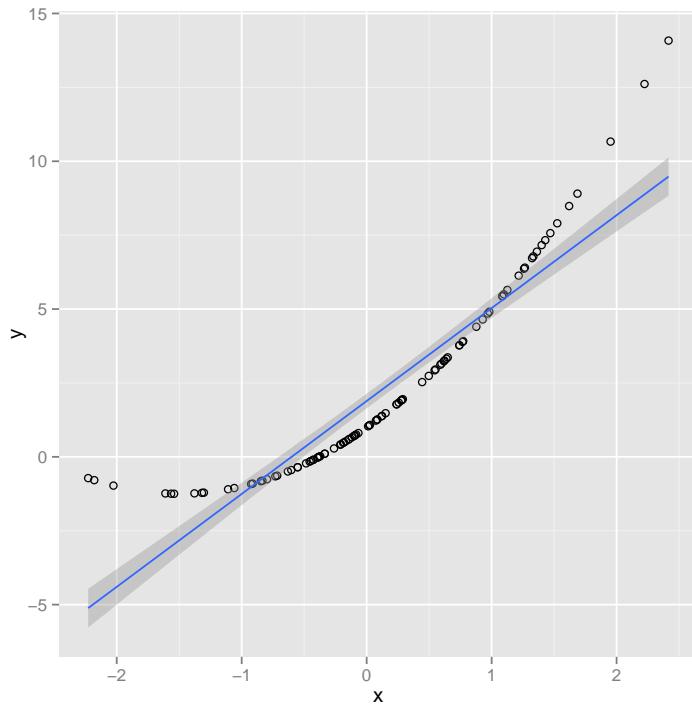
```
> ggplot(df, aes(x=value)) +
+   geom_histogram(aes(y=..density..),
+                 binwidth=.5,
+                 colour="black", fill="white") +
+   geom_density(alpha=.2, fill="#FF6666") +
+   geom_vline(aes(xintercept=mean(value, na.rm=T)),
+             color="red", linetype="dashed", size=1)
```



6.2 Line Graphs, Smoothing

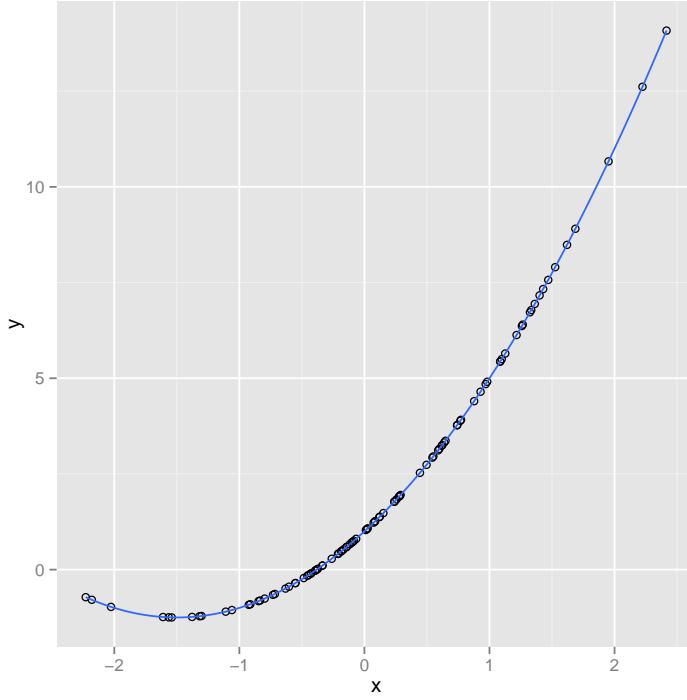
Say we have a time series and would like to plot it in ggplot. We can plot it as a scatterplot, then draw a straight line through it using a linear regression

```
> x = rnorm(100)
> y = 1 + 3*x + x^2
> df = data.frame(y = y, x = x)
> ggplot(df, aes(x=x,y=y)) +
+ geom_point(shape=1) +
+ geom_smooth(method=lm, se=T)
```



Alternativley, we could use LOESS to fit the data points, which fits a polynomial curve instead of a straight line.

```
> ggplot(df, aes(x=x,y=y)) +  
+ geom_point(shape=1) +  
+ geom_smooth()
```



6.3 Faceting the Electric Load Data

Faceting is a way to see partitions of data side by side. Let's refer back to the Kaggle Load Data. Say we would like to do a comparison across the weekdays Monday vs Tuesday vs Wednesday, etc. `ggplot` likes working with data in the long format, so we will refer to the dataset `data.long` we created earlier.

```
> head(data.long)
```

	zone_id	year	month	day	date	variable	value
1	1	2004	1	1	2004-01-01	h1	16853
2	1	2004	1	2	2004-01-02	h1	14155
3	1	2004	1	3	2004-01-03	h1	14439
4	1	2004	1	4	2004-01-04	h1	11273
5	1	2004	1	5	2004-01-05	h1	10750
6	1	2004	1	6	2004-01-06	h1	15742

First, we need more data munging. We know that date and variable together form a datetime. Again, let us create a proper datetime object.

```
> hour = as.numeric(str_extract(data.long$variable, "\d+"))
> datetime = data.long$date + 60 * 60 * hour
> data.long$datetime = datetime
> head(data.long)
```

```

zone_id year month day      date variable value      datetime
1       1 2004    1   1 2004-01-01      h1 16853 2004-01-01 01:00:00
2       1 2004    1   2 2004-01-02      h1 14155 2004-01-02 01:00:00
3       1 2004    1   3 2004-01-03      h1 14439 2004-01-03 01:00:00
4       1 2004    1   4 2004-01-04      h1 11273 2004-01-04 01:00:00
5       1 2004    1   5 2004-01-05      h1 10750 2004-01-05 01:00:00
6       1 2004    1   6 2004-01-06      h1 15742 2004-01-06 01:00:00

> str(data.long)

'data.frame': 792000 obs. of 8 variables:
$ zone_id : int 1 1 1 1 1 1 1 1 1 ...
$ year     : int 2004 2004 2004 2004 2004 2004 2004 2004 2004 ...
$ month    : int 1 1 1 1 1 1 1 1 1 ...
$ day      : int 1 2 3 4 5 6 7 8 9 10 ...
$ date     : POSIXct, format: "2004-01-01" "2004-01-02" ...
$ variable: Factor w/ 24 levels "h1","h2","h3",...: 1 1 1 1 1 1 1 1 1 ...
$ value    : num 16853 14155 14439 11273 10750 ...
$ datetime: POSIXct, format: "2004-01-01 01:00:00" "2004-01-02 01:00:00" ...

```

One thing we should make sure is that `zone_id` is recognized as a categorical variable. While it takes on values that are pure numbers, it will be helpful for us to define it as a categorical variable so we can take advantage of groups in our plots

```

> data.long$zone_id = as.factor(data.long$zone_id)
> head(data.long)

zone_id year month day      date variable value      datetime
1       1 2004    1   1 2004-01-01      h1 16853 2004-01-01 01:00:00
2       1 2004    1   2 2004-01-02      h1 14155 2004-01-02 01:00:00
3       1 2004    1   3 2004-01-03      h1 14439 2004-01-03 01:00:00
4       1 2004    1   4 2004-01-04      h1 11273 2004-01-04 01:00:00
5       1 2004    1   5 2004-01-05      h1 10750 2004-01-05 01:00:00
6       1 2004    1   6 2004-01-06      h1 15742 2004-01-06 01:00:00

> str(data.long)

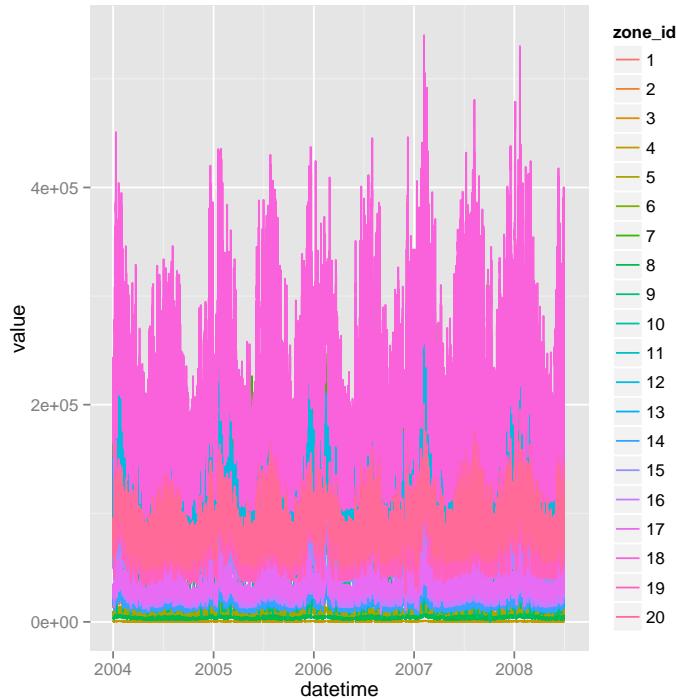
'data.frame': 792000 obs. of 8 variables:
$ zone_id : Factor w/ 20 levels "1","2","3","4",...: 1 1 1 1 1 1 1 1 1 ...
$ year     : int 2004 2004 2004 2004 2004 2004 2004 2004 2004 ...
$ month    : int 1 1 1 1 1 1 1 1 1 ...
$ day      : int 1 2 3 4 5 6 7 8 9 10 ...
$ date     : POSIXct, format: "2004-01-01" "2004-01-02" ...
$ variable: Factor w/ 24 levels "h1","h2","h3",...: 1 1 1 1 1 1 1 1 1 ...
$ value    : num 16853 14155 14439 11273 10750 ...
$ datetime: POSIXct, format: "2004-01-01 01:00:00" "2004-01-02 01:00:00" ...

```

One way to see many partitions of a dataset in a graph is to group the data points and color them according to their partitions. Alternatively, we could also generate multiple graphs and put them side by side, faceting them. For this example we will plot both.

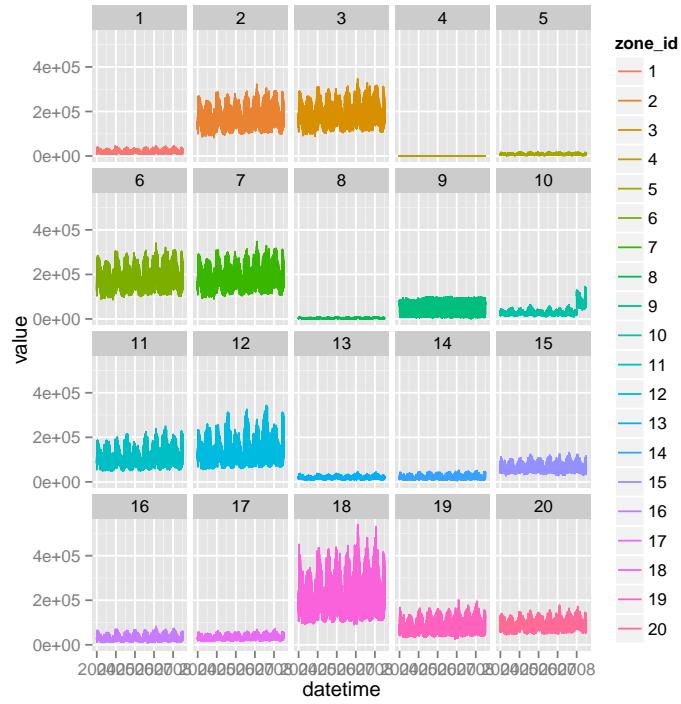
First, we define the aes by plotting datetime on the x-axis, and value on the y-axis. The color of the lines will be determined by zone_id

```
> ggplot(data.long, aes(x = datetime, y = value, color = zone_id)) +
+ geom_line()
```



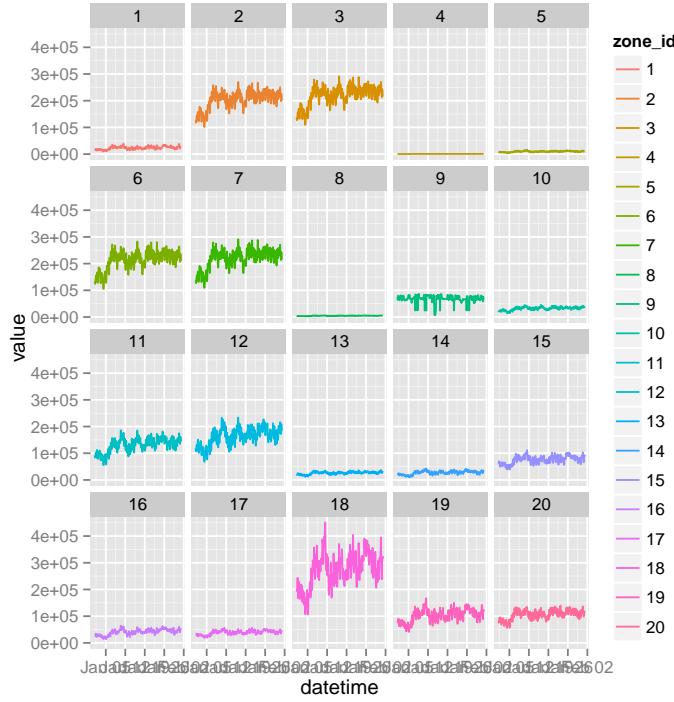
To split this one graph into a grid of smaller graphs for side by side comparison, we add facet_wrap. We will facet on zone_id so that each zone has a separate graph.

```
> ggplot(data.long, aes(x = datetime, y = value, color = zone_id)) +
+ geom_line() + facet_wrap(~zone_id, ncol = 5)
```



These graphs don't tell us much because the line graphs have too many points. Instead, we could take a subset of this data by looking at the time period January 2004 only

```
> ggplot(subset(data.long, (year == 2004 & month == 1)),
+         aes(x = datetime, y = value, color = zone_id)) +
+     geom_line() + facet_wrap(~zone_id, ncol = 5)
```



7 Data Modeling

Models are mathematical constructs that are used to explain data. What does it mean to “explain” data? In social sciences where statistics is most often used, data is never consistent. It always has variation, whether systematic variation caused by changes in a controlled environment or by unknown randomness. Models try to identify these variations, and explain why they happen. Within this field of models there are two strong schools of thought: **predictive models** and parsimonious, **data models**.

Predictive models do their utmost best to predict new data points. These models identify variation, but they do not necessarily come up with intuitive ways to explain the variation. A predictive model can be incredibly complex, but will still be popular if it makes good predictions. This school of thought is popular in machine learning.

Data models will not be as good in prediction as predictive models, but they aim to be better at explaining variance. These models are designed to be simple and parsimonious. They utilize only a few variables and the model formula has a simple form. Usually data models are appropriate when we need to make **inferences**, that is we need to understand to what degree each variable impacts the outcome, which variables are significant, what the covariances are...

7.1 Regression

Performing a linear regression is done with the lm command. Say we have data that, unknown to us, is generated using the formula $y = 1 + 3x + x^2$. If we wanted to fit a linear model to data points x and y, we could come up with the fit

```
> x = rnorm(100)
> y = 1 + 3*x + x^2 + rnorm(100)
> y.lm = lm(y ~ x)
> summary(y.lm)

Call:
lm(formula = y ~ x)

Residuals:
    Min      1Q  Median      3Q     Max 
-2.9056 -1.2193 -0.3100  0.8329  7.8503 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept)  2.0811     0.1845   11.28 <2e-16 ***
x            2.5973     0.1726   15.05 <2e-16 ***
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 1.844 on 98 degrees of freedom
Multiple R-squared:  0.698,    Adjusted R-squared:  0.6949 
F-statistic: 226.5 on 1 and 98 DF,  p-value: < 2.2e-16
```

7.2 Generalized Linear Models

7.3 Neural Networks

7.4 CART: Classification and Regression Tree

8 Advanced Data Analysis

8.1 Penalized Models

8.2 Boosting

9 Open Source Development

9.1 devtools

9.2 Roxygen

9.3 Non-R Source Code